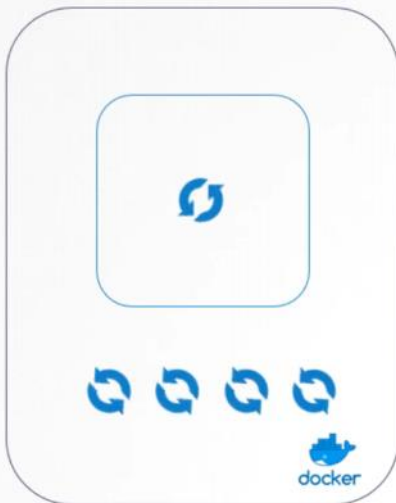


# Docker Security

## Security

```
▶ docker run ubuntu sleep 3600
```



Host



# Security



Container

```
ps aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	4528	828	?	Ss	03:06	0:00	sleep 3600

## Security - Users



Host

```
ps aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
project	3720	0.1	0.1	95500	4916	?	R	06:06	0:00	sshd: project@pts/0
project	3725	0.0	0.1	95196	4132	?	S	06:06	0:00	sshd: project@notty
project	3727	0.2	0.1	21352	5340	pts/0	Ss	06:06	0:00	-bash
root	3802	0.0	0.0	8924	3616	?	Sl	06:06	0:00	docker-containerd-shim -namespace m
root	3816	1.0	0.0	4528	828	?	Ss	06:06	0:00	sleep 3600

```
ps aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	4528	828	?	Ss	03:06	0:00	sleep 3600

```
docker run --user=1000 ubuntu sleep 3600
```

```
ps aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
1000	1	0.0	0.0	4528	828	?	Ss	03:06	0:00	sleep 3600

## Dockerfile

```
FROM ubuntu
```

```
USER 1000
```

```
▶ docker build -t my-ubuntu-image .
```

```
▶ docker run my-ubuntu-image sleep 3600
```

```
▶ ps aux
```

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
1000	1	0.0	0.0	4528	828	?	Ss	03:06	0:00	sleep 3600

# Linux Capabilities



CHOWN

DAC

KILL

SETFCAP

SETPCAP

SETGID

SETUID

NET\_BIND

NET\_RAW

MAC\_ADMIN

BROADCAST

NET\_ADMIN

SYS\_ADMIN

SYS\_CHROOT

AUDIT\_WRITE

MANY MORE

/usr/include/linux/capability.h

# Linux Capabilities



```
docker run ubuntu
```

CHOWN

DAC

KILL

SETFCAP

SETPCAP

SETGID

SETUID

NET\_BIND

NET\_RAW

MAC\_ADMIN

BROADCAST

NET\_ADMIN

SYS\_ADMIN

SYS\_CHROOT

AUDIT\_WRITE

MANY MORE

# Linux Capabilities



```
docker run --cap-add MAC_ADMIN ubuntu
```

CHOWN

DAC

KILL

SETFCAP

SETPCAP

SETGID

SETUID

NET\_BIND

NET\_RAW

MAC\_ADMIN

BROADCAST

NET\_ADMIN

SYS\_ADMIN

SYS\_CHROOT

AUDIT\_WRITE

MANY MORE

# Linux Capabilities



```
docker run --cap-drop KILL ubuntu
```

CHOWN

DAC

KILL

SETFCAP

SETPCAP

SETGID

SETUID

NET\_BIND

NET\_RAW

MAC\_ADMIN

BROADCAST

NET\_ADMIN

SYS\_ADMIN

SYS\_CHROOT

AUDIT\_WRITE

MANY MORE

# Linux Capabilities



```
docker run --privileged ubuntu
```

CHOWN

DAC

KILL

SETFCAP

SETPCAP

SETGID

SETUID

NET\_BIND

NET\_RAW

MAC\_ADMIN

BROADCAST

NET\_ADMIN

SYS\_ADMIN

SYS\_CHROOT

AUDIT\_WRITE

MANY MORE

# Container Security

```
docker run --user=1001 ubuntu sleep 3600
```

```
docker run --cap-add MAC_ADMIN ubuntu
```



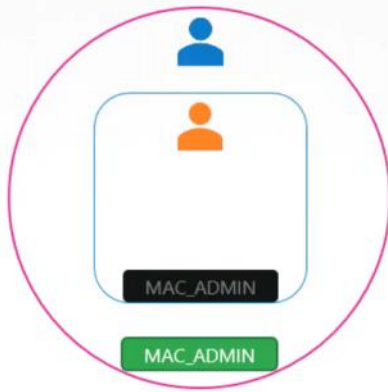
# Kubernetes Security



CONTAINER LEVEL

# Kubernetes Security

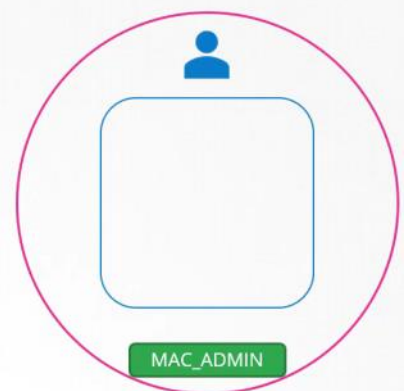




```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
```

## Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  securityContext:
    runAsUser: 1000
  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
```





# Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
      securityContext:
        runAsUser: 1000
```



# Security Context

```
apiVersion: v1
kind: Pod
metadata:
  name: web-pod
spec:
  containers:
    - name: ubuntu
      image: ubuntu
      command: ["sleep", "3600"]
      securityContext:
        runAsUser: 1000
      capabilities:
        add: ["MAC_ADMIN"]
```



**Note:** Capabilities are only supported at the container level and not at the POD level

```
controlplane ~$ kubectl exec ubuntu-sleeper -- whoami
root
```

```
apiVersion: v1
kind: Pod
metadata:
  name: multi-pod
spec:
  securityContext:
    runAsUser: 1001
  containers:
    - image: ubuntu
      name: web
      command: ["sleep", "5000"]
      securityContext:
        runAsUser: 1002
    - image: ubuntu
      name: sidecar
      command: ["sleep", "5000"]
```



So there's one at the pod level and then there's one at the container level.

So as we've learned, the container level is less security.

Context is going to override whatever is specified here.

And the sidecar container does not have a security context specified, so it's always going to use this.

```
runAsUser: 1001
```